

Deepak Mallampati

+1-330-554-0770 | deepakmallampati00@gmail.com | linkedin.com/in/deepak-mallampati | github.com/Deepak0070 |
Kent, OH

PROFESSIONAL SUMMARY

Applied AI / ML engineer shipping production healthcare AI end to end: HIPAA-conscious RAG (~35% retrieval precision, >90% grounded alignment), predictive ML pipelines (15-20% recall on high-risk patient categories at >90% precision), agentic LLM workflows (>30% hallucination reduction), and FastAPI + Docker packaging (~30% post-deploy defect reduction). Production computer vision via YOLOv8 (~30% inference latency reduction). Master's in Computer Science, Kent State. Targeting Percepta's Applied AI Engineer seat (NYC, GC-backed vertical AI for healthcare/manufacturing/energy) - production healthcare RAG with grounded answers and clinical SME stakeholder discipline (1:1 with Mosaic's "rapidly deploy agentic workflows" charter), Manga Lens 4-provider AI vision integration shipped solo on Chrome Web Store with cost-aware payload routing (proves customer-deployed multi-provider LLM orchestration), Agentic Healthcare Claims with schema-validated JSON contracts between five agents (matches Percepta's production-grade AI agents with audit-ready reasoning), and E-Farming AgriTech marketplace as solo founder (extreme-ownership and 0 → 1 product instinct). NYC relocation open, F-1 OPT US-base.

CORE COMPETENCIES

Applied AI Engineer • Production AI Agents / Workflows • RAG / Embeddings / Vector Search • Multi-Provider LLM API Integration (OpenAI, Anthropic, Google, Ollama) • LangChain / Agentic / Schema Contracts • Python / TypeScript • FastAPI / Flask / Docker / REST • CI/CD (Jenkins, GitHub Actions) • Healthcare AI / HIPAA Governance • Founder / 0-to-1 Ownership • NYC Relocation Open (F-1 OPT)

PROFESSIONAL EXPERIENCE

Progress Solutions Inc.

Jul 2025 - Present

Jr. AI/ML Engineer Trainee - Applied AI, Healthcare

USA - Healthcare Technology, Applied AI

- Built **Retrieval-Augmented Generation (RAG)** for clinical knowledge retrieval and healthcare documentation search; recursive semantic chunking + transformer embeddings improved retrieval precision ~**35%**; retrieval-grounded response alignment >**90%**.
- Developed **agentic LLM workflows** for multi-step healthcare queries (eligibility checks, care workflow navigation, documentation clarification) with structured reasoning, tool discipline, and grounding rules; agent response stability ~**25%**; hallucinations >**30%**.
- Shipped **predictive ML pipelines** (scikit-learn, XGBoost) for patient no-show, care engagement scoring, support prioritization; recall +**15-20%** on high-risk categories at >**90%** precision via class weighting, stratified sampling, threshold calibration.
- Packaged ML/LLM inference as **FastAPI / Flask** services on **Docker** with structured logging and load simulation; post-deploy defects ~**30%**.
- EHR + appointment + ticket preprocessing (Pandas, NumPy); dataset reliability >**98%**, downstream instability **-40%**.
- Authored **HIPAA-conscious data governance** and stakeholder-facing system-limitation docs; ran weekly stakeholder calls with clinical SMEs.

Energy Solutions International (Emerson)

Jun 2022 - Apr 2023

Database & DevOps Performance Engineer (Intern)

Hyderabad, India - Oil & Gas, Enterprise ERP

- Optimized SQL and **T-SQL stored procedures** on the Synthesis Order-to-Cash platform; query execution time ~**20%**, data retrieval latency ~**25%**.
- Built **CI/CD pipelines with Jenkins** for schema updates and stored procedure deployments; deployment errors >**30%**, release cycle ~**35-40%**.
- Designed **SQL Server DMVs and Grafana** performance dashboards for long-running queries, deadlocks, CPU/IO contention; incident recurrence ~**25%**, RCA speed ~**30%**.
- Implemented **RBAC and audit logging for financial modules** in a compliance-sensitive oil & gas environment.

Machine Learning Engineer

Hyderabad, India - Nonprofit, Video Analytics, Applied NLP

- Built **transformer-based video summarization** over 5,000+ recorded sessions; hierarchical summarization + timestamp-aligned clip extraction cut manual review **60-70%**.
- Implemented **document Q&A with RAG**; hallucinations **~30%**, contextual accuracy **>85%**.
- Deployed as **Flask** API with a lightweight web interface for non-technical staff.

PROJECTS

Manga Lens - Multi-Provider AI Vision Chrome Extension (shipped, Chrome Web Store)

Real-time manga / webtoon translation extension shipped to the Chrome Web Store. **Multi-provider abstraction layer** integrating four AI vision providers (Claude Sonnet, GPT-4o-mini, GPT-4.1-Nano, Ollama/minicpm-v) with per-provider payload handling (WebP for cloud, JPEG for local), cost-aware routing, circuit-breaker fallback. Manifest V3 service workers; multi-section capture pipeline for tall panels with coordinate remapping; 7-day cache; per-domain selector configs for 29 manga/webtoon sites; narrowed host permissions and privacy policy.

TypeScript, Manifest V3, Service Workers, OpenAI / Anthropic / Google / Ollama, multi-provider integration

Agentic Healthcare Claims Processing & Fraud Risk Intelligence System

Five-stage multi-agent pipeline (Intake Normalization, Validation, Consistency Analysis, Duplicate Detection, Fraud Risk Scoring) with **schema-validated JSON contracts between agents** to prevent cascading hallucinations. RAG-grounded CPT/ICD validation against vector-indexed policies. Audit-ready explainable risk scoring with reasoning traces.

Python, LangChain, FastAPI, vector search, GPT-4, schema contracts

Dream Decoder - Multimodal Creative Intelligence Platform

Full-stack FastAPI + React/TypeScript/Vite/Tailwind app interpreting dreams, generating poetic reinterpretations, and synthesizing dreamscapes via coordinated multimodal APIs (Perplexity Sonar for interpretation, Replicate for 16:9 visuals). Introduced **intermediate structured prompt transformation layers** - improved contextual alignment **~30%** and first-pass image success **~25-30%** over naive direct prompting.

FastAPI, React, TypeScript, Vite, Tailwind, Perplexity Sonar, Replicate

Driver Drowsiness Detection - Real-Time YOLOv8 Fatigue Monitoring

Replaced a two-stage CNN pipeline with a **unified YOLOv8** detection-and-classification model; reduced inference latency **~30%**. Sliding-window confidence aggregation suppressed blink-driven false positives **~25%**; adaptive frame skipping; NMS tuning for stable real-time operation via OpenCV.

YOLOv8, OpenCV, PyTorch, real-time inference

E-Farming Digital Marketplace (Founder & Full-Stack Developer)

Independently designed and shipped a PHP/MySQL/Bootstrap AgriTech marketplace connecting small-scale farmers with buyers. Onboarded **80-120 active users** in phase one; cart, reviews, and community blogging features to drive repeat engagement.

PHP, MySQL, Bootstrap, full-stack, founder-led

EDUCATION

Master of Science, Computer Science - Kent State University

Aug 2023 - May 2025

Focus: Applied Machine Learning, NLP, Database Systems.

Bachelor of Technology, Computer Science - KL University

Jun 2019 - May 2023

Founder / Lead of E-Farming platform (university project).

CERTIFICATIONS

Deep Learning Specialization - DeepLearning.AI (Coursera)

2024

SKILLS

LLM / GenAI: RAG (chunking, embeddings, vector search), agentic workflows, multi-step agents, schema-validated JSON contracts, structured outputs, evaluation pipelines, guardrails, grounding, LangChain, multi-provider AI integration (OpenAI, Anthropic, Google, Ollama)

Languages: Python (primary), TypeScript, JavaScript, SQL (T-SQL, PostgreSQL, SQLite), C++

ML Frameworks: PyTorch, scikit-learn, XGBoost, Hugging Face Transformers, Diffusers, LoRA fine-tuning

Backend / Inference: FastAPI, Flask, REST, Docker, structured logging, load simulation, microservices

Frontend: React, TypeScript, Vite, Tailwind, Manifest V3 / Chrome Extensions

DevOps: Docker, Jenkins, GitHub Actions, Grafana, CI/CD, cloud-ready deployment

Data Governance: HIPAA-conscious de-identification, data lineage, evaluation audit trails, RBAC, audit logging

Domains: Vertical AI (Healthcare / Manufacturing-adjacent / Energy-adjacent), Production AI Agents, Founder / 0-to-1